



# Control Taskflow Graph (CTFG) Programming Model

---

Dr. Tsung-Wei (TW) Huang

Department of Electrical and Computer Engineering  
University of Wisconsin at Madison, Madison, WI

<https://taskflow.github.io/>





# Takeaways

---


- Learn the control Taskflow graph (CTFG) programming model
- Compare CTFG with existing models and recognize their limitations
- Understand the scheduling flow of a CTFG
- Showcase some real-world applications of CTFG
- Conclude the talk




# Control Taskflow Graph (CTFG) Programming Model

- A key innovation that distinguishes Taskflow from existing models

```
auto [init, cond, yes, no] =
taskflow.emplace(
[] () { },
[] () { return 0; },
[] () { std::cout << "yes"; },
[] () { std::cout << "no"; }
);
cond.succeed(init)
    .precede(yes, no);
```



```
auto [init, opt, cond, stop] =
taskflow.emplace(
[&](){ initialize_data_structure(); },
[&](){ some_optimizer(); },
[&](){ return converged() ? 1 : 0; },
[&](){ std::cout << "done!\n"; }
);
opt.succeed(init).precede(cond);
converged.precede(opt, stop);
```



CTFG goes beyond the limitation of traditional DAG-based frameworks (no in-graph control flow).




# Power of CTFG Programming Model

- Enables very efficient overlap among tasks and control flow

```
auto cond_1 = taskflow.emplace([](){ return run_B() ? 0 : 1; });
auto cond_2 = taskflow.emplace([](){ return run_G() ? 0 : 1; });
auto cond_3 = taskflow.emplace([](){ return loop() ? 0 : 1; });
cond_1.precede(B, E);
cond_2.precede(G, H);
cond_3.precede(cond_3, L);
```

This type of end-to-end parallelism is almost impossible to achieve using existing task-parallel programming (TPP) libraries.






# Condition Task

- A callable that returns an integer indicating which successor task to run
  - The returned integer represents the index of the successor task in the dependency list

```
tf::Taskflow taskflow;
auto [A, B, C, D, E] = taskflow.emplace(
    [&]()
    { std::cout << "A\n"; },
    [&() -> int {
        std::cout << "B is a condition task\n";
        return 2;
    }, // B is a condition task
    [&() { std::cout << "C\n"; },
    [&() { std::cout << "D\n"; },
    [&() { std::cout << "E\n"; }

);
A.precede(B);
B.precede(C, D, E);
```



If the returned integer index is beyond the valid successor range (e.g., return 100), the scheduler will not insert any task to run.




# Iterative Control Flow

- Use condition tasks to build an iterative control flow within a task graph

```
tf::Taskflow taskflow;

int i;
auto [init, cond, body, back, done] = taskflow.emplace(
    [&]{ std::cout << "i=0"; i=0; },
    [&]{ std::cout << "while i<5\n"; return i < 5 ? 0 : 1; },
    [&]{ std::cout << "i++=\n" << i++ << '\n'; },
    [&]{ std::cout << "back\n"; return 0; },
    [&]{ std::cout << "done\n"; }
);
init.precede(cond);           // init
cond.precede(body, done);    // while i<5
body.precede(back);          // i++
back.precede(cond);          // back
```


The cond task acts as the *while condition*, checking if *i* is less than 5.



# Switch-Case Control Flow

- Use condition tasks to build a switch-case control flow within a task graph

```
auto [source, swcond, case1, case2, case3, target] = taskflow.emplace(  
    [](){ std::cout << "source\n"; },  
    [](){ std::cout << "switch\n"; return rand()%3; },  
    [](){ std::cout << "case 1\n"; return 0; },  
    [](){ std::cout << "case 2\n"; return 0; },  
    [](){ std::cout << "case 3\n"; return 0; },  
    [](){ std::cout << "target\n"; }  
);  
  
source.precede(swcond);  
swcond.precede(case1, case2, case3);  
target.succeed(case1, case2, case3);
```



The three case tasks (case 1, case 2, case 3) must also be a condition task because target depends only on one case task.




# Condition Tasks can Handle Complicated Scenarios

- Use condition tasks to build a non-deterministic control flow

```
auto A = taskflow.emplace([&](){ });
auto B = taskflow.emplace([&](){ return rand()%2; } );
auto C = taskflow.emplace([&](){ return rand()%2; } );
auto D = taskflow.emplace([&](){ return rand()%2; } );
auto E = taskflow.emplace([&](){ return rand()%2; } );
auto F = taskflow.emplace([&](){ return rand()%2; } );
auto G = taskflow.emplace([&]());
```

```
A.precede(B).name("init");
B.precede(C, B).name("flip-coin-1");
C.precede(D, B).name("flip-coin-2");
D.precede(E, B).name("flip-coin-3");
E.precede(F, B).name("flip-coin-4");
F.precede(G, B).name("flip-coin-5");
G.name("end");
```

Each condition task flips a binary coin to decide the next task to run:  
 0: goes to the next task  
 1: goes back to the original task



# Multi-condition Task

- A generalized version of condition task to jump to multiple successor tasks

```
tf::Executor executor;  
tf::Taskflow taskflow;
```


```
auto A = taskflow.emplace([&]() -> tf::SmallVector<int> {  
    std::cout << "A\n";  
    return {0, 2};  
});
```

```
auto B = taskflow.emplace([&](){ std::cout << "B\n"; });  
auto C = taskflow.emplace([&](){ std::cout << "C\n"; });  
auto D = taskflow.emplace([&](){ std::cout << "D\n"; });
```

```
A.precede(B, C, D);
```

```
executor.run(taskflow).wait();
```

tf::SmallVector<int> is similar to std::vector but comes with small buffer optimization.



If the returned integer index is beyond valid successor range (e.g., return -1), the scheduler will not insert any task to run.





# Takeaways

---

- Learn the control Taskflow graph (CTFG) programming model
- Compare CTFG with existing models and recognize their limitations
- Understand the scheduling flow of a CTFG
- Showcase some real-world applications of CTFG
- Conclude the talk




# How Do Existing TPP Libraries Handle Control Flow?

- Most existing TPP libraries rely on **Directed Acyclic Graph (DAG) models**
  - Do not natively support conditional execution or iterative execution (e.g., cycle)
  - As a result, all the control-flow decision making must happen outside the task graph
- A common approach or workaround is as follows:
  - You manually partition the task graph around control-flow points – *out-of-graph control flow*
  - Then, each of those partitioned graphs synchronizes at the control-flow boundaries
  - This approach works just fine but not optimal, as it limits **end-to-end parallelism** and increases user-side code complexity

```
tf::Taskflow G;
auto X = G.emplace([](){});  
auto Y = G.emplace([](){  
    return converged() ? 1 : 0;  
});  
Y.precede(Y, Z).succeed(X);
executor.run(G).wait();
```

Manually partition the task graph using a do-while loop (out-of-graph control flow)



Iterative synchronization cost at the end of each iteration.

```
tbb::flow::graph1 X, Y;  
tbb::flow::graph Z;  
X.run(); // sync  
do {  
    Y.run(); // sync  
} while (!converged());  
Z.run();
```






# Performance: In-graph vs Out-of-graph Control Flow

```
// in-graph control flow (CTFG)
tf::Taskflow G;
auto X = G.emplace([](){});
auto Y = G.emplace([](){
    return converged() ? 1 : 0;
});
Y.precede(Y, Z).succeed(X);
executor.run(G).wait();
```

```
// out-of-graph control flow
tf::Taskflow X, Y, Z;
executor.run(X).wait(); // sync
do {
    executor.run(Y).wait(); // sync
} while (!converged());
executor.run(Z).wait();
```




# How to Partition this Task Graph?

- Three condition tasks (`cond_1`, `cond_2`, `cond_3`) within a task graph

```
auto cond_1 = taskflow.emplace([](){ return run_B() ? 0 : 1; });
auto cond_2 = taskflow.emplace([](){ return run_G() ? 0 : 1; });
auto cond_3 = taskflow.emplace([](){ return loop() ? 0 : 1; });

cond_1.precede(B, E);
cond_2.precede(G, H);
cond_3.precede(cond_3, L);
```

Again, without CTFG, this type of end-to-end parallelism is very difficult to achieve using existing task-parallel programming libraries.





# Takeaways


---

- Learn the control Taskflow graph (CTFG) programming model
- Compare CTFG with existing models and recognize their limitations
- Understand the scheduling flow of a CTFG
- Showcase some real-world applications of CTFG
- Conclude the talk






# Scheduling Flow of a Control Taskflow Graph





# Example: In-graph Iterative Control flow

- 1. Starts with init because it has zero strong and weak dependencies**
  - If your CTFG doesn't have any zero-dependency tasks, it won't be schedulable
- 2. Moves on to the condition task cond**
  - If cond returns 0, the scheduler enqueues cond and runs it again
  - If cond returns 1, the scheduler enqueues stop and then moves on



Taskflow always starts with tasks that have zero strong and weak dependencies (i.e., init in this example).

task	# strong dep	# weak dep
init	0	0
cond	1	1
stop	0	1







# Conditional Tasking can be Easy to Make Mistakes!

- It is your responsibility to ensure the given CTFG is valid
  - A valid CTFG is a task graph that is schedulable and race-free (no task race)
- Some common pitfalls in CTFG programming below:



Not schedulable  
(#1 mistake in CTFG)



Task race occurs on D when C returns 0



P and M must not return 1 and 0





# Takeaways


---

- Learn the control Taskflow graph (CTFG) programming model
- Compare CTFG with existing models and recognize their limitations
- Understand the scheduling flow of a CTFG
- **Showcase some real-world applications of CTFG**
- Conclude the talk




# Use Case #1: K-means Clustering

- Implemented the iterative centroid update using CTFG<sup>1</sup>



# Use Case #2: VLSI Detailed Placement

- Implemented an iterative gate location update algorithm using CTFG<sup>1</sup>



# Use Case #3: Robotics Processing Planning



## Graph Taskflow Generator Revision #17

Merged Levi-Armstrong merged 7 commits into `tesseract-robotics:master` from `marip8:update/graph_taskf...`

Conversation 28 Commits 7 Checks 0 Files changed 5

marip8 commented on Jan 19, 2021 · edited

This PR revises the graph Taskflow generator class to be easier to use.

### Assumptions

- End conditions
  - The required error task is only applicable to conditional nodes. I assign the error task as the connection at index 0 to all conditional tasks
  - The done callback is applicable to both conditional and non-conditional nodes.
    - If a conditional node has any edges, I do not connect that node to the done callback
    - If a conditional node is a leaf node (i.e. has no defined edges), the connection at index 0 is to the error task (as mentioned above) and the connection at index 1 is to the done task
    - If a non-conditional node is a leaf node, its only connection is to the done callback
- No graph checking is implemented for loops and "island" states. I think we should allow users the freedom to define their graphs in whatever way possible. It's easy enough to dump the taskflow to a file and understand if the graph looks correct

All of these assumptions should be documented in the header file





# Takeaways

---

- Learn the control Taskflow graph (CTFG) programming model
- Compare CTFG with existing models and recognize their limitations
- Understand the scheduling flow of a CTFG
- Showcase some real-world applications of CTFG
- Conclude the talk





# Question?

## Static Task Graph Programming (STGP)

// Live: <https://godbolt.org/z/j8hx3xnnx>

```
tf::Taskflow taskflow;
tf::Executor executor;
auto [A, B, C, D] = taskflow.emplace(
    [](){ std::cout << "TaskA\n"; },
    [](){ std::cout << "TaskB\n"; },
    [](){ std::cout << "TaskC\n"; },
    [](){ std::cout << "TaskD\n"; });
A.precede(B, C);
D.succeed(B, C);
executor.run(taskflow).wait();
```



Taskflow: <https://taskflow.github.io>

## Dynamic Task Graph Programming (DTGP)

// Live: <https://godbolt.org/z/T87PrTdx>

```
tf::Executor executor;
auto A = executor.silent_dependent_async[]{
    std::cout << "TaskA\n";
};
auto B = executor.silent_dependent_async[]{
    std::cout << "TaskB\n";
}, A;
auto C = executor.silent_dependent_async[]{
    std::cout << "TaskC\n";
}, A;
auto D = executor.silent_dependent_async[]{
    std::cout << "TaskD\n";
}, B, C;
executor.wait_for_all();
```

