



# Taskflow: A General-purpose Task-parallel Programming System

---

Dr. Tsung-Wei (TW) Huang

Department of Electrical and Computer Engineering  
University of Wisconsin at Madison, Madison, WI

<https://taskflow.github.io/>





# Takeaways


---

- Understand the importance of task-parallel programming
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Showcase real-world applications of Taskflow
- Conclude the talk



# Why Parallel Computing?

- Advances performance to a new level previously out of reach






# Modern Hardware is Designed to Run in Parallel

- **Intel Haswell microarchitecture**

- Released in June 2013
- Typically comes with four cores
- Has an integrated GPU
- 1.4 B transistors with 22 nm technology
- Sophisticated design for ILP acceleration
- Deep pipeline – 16 stages




- **Superscalar architecture**

- Can issue and complete multiple independent instructions per cycle

- **Supports hyper-threading tech (HTT)**


- Allows a single physical CPU core to appear as two logical processors to the OS

If you don't do parallel programming, you are not utilizing your hardware efficiently ...




# Today's Parallel Computing Problem is Very Irregular

- Computational task graph of a GPU-parallel circuit simulation workload<sup>1</sup>
  - > 500M gates and nets
  - > 1000 async kernel tasks
  - > 1000 dependencies
  - > hours to finish




# Another Example of Irregular Parallel Workload

- CPU-parallel VLSI static timing analysis algorithm



Analysis algorithm is modeled as a task graph

> 100M tasks  
> 150M dependencies






# Parallelizing such Irregular Workloads is Not Easy ...

- You need to deal with A LOT OF technical details

- Parallelism abstraction (software + hardware)
- Concurrency control
- Synchronization
- Task and data race avoidance
- Dependency constraints
- Scheduling efficiencies (load balancing)
- Programming productivity
- Performance portability
- ...



- And, don't forget about trade-offs

- Performance vs Developer's intent


We want a solution that can sit on top to help programmers manage these details as much as possible because programmers care how fast (performance + productivity) they can get things done!



# Why Task-parallel Programming (TPP)?

- TPP is an effective solution for parallelizing irregular workloads
  - Captures developers' intention in decomposing an algorithm into a *top-down* task graph
  - Delegates difficult scheduling details (e.g., load balancing) to an optimized runtime
- Modern parallel programming libraries are moving towards task parallelism
  - OpenMP 4.0 task dependency clauses (`omp depend`)
  - C++26 execution control library (`std::exec`)
  - TBB flow graph (`tbb::flow::graph`)
  - Taskflow control Taskflow graph (CTFG) model
  - ... (many others)

Task graph encapsulates functions as tasks and their dependencies as edges.




**OpenMP**

**StarPU**

**kokkos**



**PaRSEC**





# Takeaways

---


- Understand the importance of task-parallel programming
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Showcase real-world applications of Taskflow
- Conclude the talk



# Static Task Graph Programming in Taskflow<sup>1</sup>

```
#include <taskflow/taskflow.hpp> // Live: https://godbolt.org/z/j8hx3xnnx
```


```
int main(){
    tf::Taskflow taskflow;
    tf::Executor executor;
    auto [A, B, C, D] = taskflow.emplace(
        [] () { std::cout << "TaskA\n"; },
        [] () { std::cout << "TaskB\n"; },
        [] () { std::cout << "TaskC\n"; },
        [] () { std::cout << "TaskD\n"; }
    );
    A.precede(B, C);
    D.succeed(B, C);
    executor.run(taskflow).wait();
    return 0;
}
```




# Control Taskflow Graph (CTFG) Programming Model

- A key innovation that distinguishes Taskflow from existing libraries

```
auto [init, cond, yes, no] =
taskflow.emplace(
    [] () { },
    [] () { return 0; },
    [] () { std::cout << "yes"; },
    [] () { std::cout << "no"; }
);
cond.succeed(init)
    .precede(yes, no);
```



```
auto [init, opt, cond, stop] =
taskflow.emplace(
    [&]() { initialize_data_structure(); },
    [&]() { some_optimizer(); },
    [&]() { return converged() ? 1 : 0; },
    [&]() { std::cout << "done!\n"; }
);
opt.succeed(init).precede(cond);
converged.precede(opt, stop);
```



CTFG goes beyond the limitation of traditional DAG-based frameworks (no in-graph control flow).




# Power of CTFG Programming Model

- Enables very efficient overlap among tasks alongside control flow

```
auto cond_1 = taskflow.emplace([](){ return run_B() ? 0 : 1; });
auto cond_2 = taskflow.emplace([](){ return run_G() ? 0 : 1; });
auto cond_3 = taskflow.emplace([](){ return loop() ? 0 : 1; });

cond_1.precede(B, E);
cond_2.precede(G, H);
cond_3.precede(cond_3, L);
```

This type of parallelism is almost impossible to achieve using existing TPP libraries.





# Takeaways


---

- Understand the importance of task-parallel programming
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Showcase real-world applications of Taskflow
- Conclude the talk



# Why STGP Alone is Not Sufficient?

- In STGP, the graph structure must be known up front before execution
  - Execution of STGP is based on the *construct-and-run* model
- Lack of overlap between task graph construction and task graph execution
  - For large task graphs, this overlap can sometimes bring a significant performance advantage
- Lack of flexibility for dynamically expressing task graph parallelism
  - Task graph structure cannot depend on runtime values or control-flow results



# Dynamic Task Graph Programming in Taskflow



// Live: <https://godbolt.org/z/j76ThGbWK>

```
tf::Executor executor;

auto A = executor.silent_dependent_async([](){
    std::cout << "TaskA\n";
});

auto B = executor.silent_dependent_async([](){
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async([](){
    std::cout << "TaskC\n";
}, A);
auto [D, Fu] = executor.dependent_async([](){
    std::cout << "TaskD\n";
}, B, C);

Fu.wait();
```



Specify arbitrary task dependencies on previously created tasks





# DTGP is Flexible for Runtime-driven Execution

- Assemble task graphs driven by runtime variables and control-flow results

```
if (a == true) {  
    G1 = build_task_graph1();  
    if (b == true) {  
        G2 = build_task_graph2();  
        G1.precede(G2);  
        if (c == true) {  
            ... // defined other TGP  
        }  
    }  
    else {  
        G3 = build_task_graph3();  
        G1.precede(G3);  
    }  
}
```

```
G1 = build_task_graph1();  
G2 = build_task_graph2();  
if (G1.num_tasks() == 100) {  
    G1.precede(G2);  
}  
else {  
    G3 = build_task_graph3();  
    G1.precede(G2, G3);  
    if(G2.num_dependencies()>=10){  
        ... // define another TGP  
    } else {  
        ... // define another TGP  
    }  
}
```

This type of dynamic task graph is very difficult to achieve using static task graph programming ...





# Takeaways

---


- Understand the importance of task-parallel programming
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- **Showcase real-world applications of Taskflow**
- Conclude the talk




# VLSI Static Timing Analysis

- **Task-parallel timing propagation<sup>1</sup>**

- Task: per-pin propagation function
  - Ex: cell delay, net delay calculator
- Edge: pin-to-pin dependency
  - Ex: intra-/inter-gate dependencies



↓ Derive a timing propagation task graph





```
ot> report_timing # report the most critical path
Startpoint    : inp1
Endpoint      : f1:D
Analysis type : min
-----
Type Delay Time Dir Description
-----
port 0.000 0.000 fall inp1
pin  0.000 0.000 fall u1:A (NAND2X1)
...
pin  0.000 2.967 fall f1:D (DFFNEGX1)
...
slack -23.551 VIOLATED
```

↑ Evaluate and report violated data paths

# How Good is Task-parallel STA?

- **OpenTimer v1: levelization-based (or loop-parallel) timing propagation<sup>1</sup>**
  - Implemented using OpenMP “parallel\_for” primitive
- **OpenTimer v2: task-parallel timing propagation<sup>2</sup>**
  - Implemented using Taskflow STGP



 Task-parallelism allows us to more asynchronously parallelize the timing propagation




<sup>1</sup>: Tsung-Wei Huang and Martin Wong, “OpenTimer: A High-Performance Timing Analysis Tool,” *IEEE/ACM ICCAD*, 2015

<sup>2</sup>: Tsung-Wei Huang, et al, “OpenTimer v2: A New Parallel Incremental Timing Analysis Engine,” *IEEE TCAD*, 2022

# Other Industrial Applications of Taskflow

- At the time of this talk, Taskflow has weekly downloads of 7-9K
  - Quantum computing (e.g., Xanadu uses Taskflow in JET to simulate quantum tensor network)
  - Bioinformatics (e.g., Roche uses Taskflow to parallelize its DNA sequencing software)
  - Computer graphics (e.g., Vulkan recommends Taskflow for parallelizing rendering engines)
  - Embedded systems (e.g., Tesseract uses Taskflow to design robotic planning software)
  - Scientific computing (e.g., Deal.II uses Taskflow to parallelize finite element analysis)
  - ...





# Takeaways

---

- Understand the importance of task-parallel programming
- Program static task graph parallelism using Taskflow
- Program dynamic task graph parallelism using Taskflow
- Showcase real-world applications of Taskflow
- Conclude the talk





# Using Taskflow is EXTREMELY EASY

- **Taskflow is a header-only library built entirely with standard C++ libraries**
  - No wrangling with installation – just copy the headers and tell your compiler where to find them

```
# clone the Taskflow project
```

```
~$ git clone https://github.com/taskflow/taskflow.git
~$ cd taskflow
```

```
# compile your program and tell your compiler where to find Taskflow header files
```

```
~$ g++ -std=c++20 examples/simple.cpp -I ./ -O2 -pthread -o simple
~$ ./simple
```

```
TaskA
```

```
TaskC
```

```
TaskB
```

```
TaskD
```

- **Taskflow has been evolving over the years to a stable programming system**


- Started in 2018 as a DARPA-sponsored research project to parallelize critical EDA applications





# Visualize the Execution of Your Taskflow Programs

```
# run your program with the env variable TF_ENABLE_PROFILER enabled and  
# paste the output JSON content to https://taskflow.github.io/tfprof/  
~$ TF_ENABLE_PROFILER=simple.json ./simple
```





Thank you for using Taskflow! <https://taskflow.github.io/>

---



...





# Question?



Taskflow: <https://taskflow.github.io>

## Static Task Graph Programming (STGP)

// Live: <https://godbolt.org/z/j8hx3xnnx>

```
tf::Taskflow taskflow;
tf::Executor executor;
auto [A, B, C, D] = taskflow.emplace(
    [](){ std::cout << "TaskA\n"; },
    [](){ std::cout << "TaskB\n"; },
    [](){ std::cout << "TaskC\n"; },
    [](){ std::cout << "TaskD\n"; });
A.precede(B, C);
D.succeed(B, C);
executor.run(taskflow).wait();
```



## Dynamic Task Graph Programming (DTGP)

// Live: <https://godbolt.org/z/T87PrTax>

```
tf::Executor executor;
auto A = executor.silent_dependent_async[]{
    std::cout << "TaskA\n";
};
auto B = executor.silent_dependent_async[]{
    std::cout << "TaskB\n";
}, A);
auto C = executor.silent_dependent_async[]{
    std::cout << "TaskC\n";
}, A);
auto D = executor.silent_dependent_async[]{
    std::cout << "TaskD\n";
}, B, C);
executor.wait_for_all();
```

